
הח"כ ה-121

Steal / Never Sheet

Research findings, current-state audit and structural decisions for the UX and game-design redesign. Version 8.

Steal / Never Sheet — The 121st MK · v8

Research output + current-state audit + structural decisions · Aug 2026 · Lion Blau

Grounded in reviewed captures and the live build only. End-screen mechanics and question sheet read 27 Aug. Principles layer v7. Arc revised and `data.js` read v8 (29 Aug). **SET** confirmed · **OPEN** open · **PROPOSAL** proposal awaiting sign-off · **NOW** has a clock on it Owners: **Tamar/NGO** = content, editorial, political framing · **Roman** = data, logic, backend, architecture

Scope note: UX / features / interaction. Visual direction parked in §13.

0. Current state — what is actually built (verified 27 Aug)

Verified against the live build at `romangash-star.github.io/hachach-121` plus screen captures. Everything here is observed, not inferred.

0.1 The round already has the arc's skeleton

Step	What happens	Reveal	Coins
1	<code>אמת או שקר?</code> claim card — swipe or tap	Immediate: verdict, explanation, <code>מקור: mako</code>	+25 if correct
2	<code>איך היית מצביע?</code> — the player's own position (בעד/נגד/נמנע)	None	0 — unscored, unused
3	MK predictions — <code>מה לדעתך הוא/היא הצביע?</code>	Batched via <code>חשוף את תוצאות ההצבעה</code>	+25 per correct MK
—	Topic complete (<code>כל הכבוד!</code>)	—	+100

A 3-segment progress bar in the round header already tracks these steps. **The arc is a re-sequencing, not a rebuild** — components exist, sequencing is the work.

Revised estimate: ~2-2.5 days, down from 2.5-3.5.

0.2 Structure, as built

- 8 topics × 2 סוגיות = 16 rounds, plus **bonus** סוגיות per topic (`סוגיית בונוס בנושא הזה`)
on the topic-complete screen; map at 8/8 reads `סיימת את המפה! יש עוד סוגיות בונוס`)
- Progress is counted **per topic**: 0/8 ... 8/8 . The 16 sub-rounds are invisible on the map
- Map is **free choice** — no locking, no prerequisites
- Completed topic cards go **saturated colour + green** + `הושלם` (each topic has its own hue — red, yellow, pink, blue, green)
- End flow: last reveal → `סיימת את המפה — לחגיגה!` → celebration → coin allocation → share
- Bottom nav: `מפה / דירוג / ממשלה / אודות`

0.3 The coin economy, as observed

Event	Award	Notes
אמת או שקר correct	+25	
Each correct MK prediction	+25	3 of 5 correct → +75
Topic completed	+100	סה"כ 1750 מטבעות · +100
Own vote (איך היית מצביע)	0	The most interesting input in the game earns nothing
Observed full-game total	1,750	8/8 topics

PROPOSAL The economy has no sink — and there is no allocation mechanic at all. (Verified 27 Aug: the end screen is a single free-text field, not an allocation UI.)

The player accumulates 1,750 coins across sixteen rounds. The celebration screen promises הגיע הזמן להחליט מה חשוב לכם and then presents **one free-text input**, whose contents are pasted verbatim into the share string:

המגדלור: # [free text] צברתי 1750 מטבעות, ובחרתי:] " החכ ה-121 " שיחקתי

The currency does nothing, ever. The 1750 מטבעות button sits directly above the field and is decorative. Coins are never spent, converted, unlocked against, or referenced again.

The 0/100 in the DOM is not reachable in play. Either an unbuilt allocation UI sitting in the markup, or dead code. → **Roman**. This matters: if a real mechanic is half-built, the job is *finishing* one, not *inventing* one.

The placeholder is promising a mechanic the input doesn't have. The field reads חינוך, בריאות, סביבה, ביטחון, which looks like a multi-select list of causes. Both Lion and Claude initially read this screen as an allocation UI because of it. That's a UX finding in its own right: the placeholder sets an expectation the input cannot meet.

PROPOSAL 0.3b Proposed fix — allocate across the eight topics

This isn't replacing a working mechanic; it's **filling a hole the build already has**. The screen currently makes a promise and doesn't keep it.

Peak-End Rule (Kahneman): experiences are remembered by their peak and their *ending*, not their average. The current ending is a free-text box above a dead coin counter — and per peak-end, players will remember the whole game partly through that screen. **The end-game is memory infrastructure, not a nice-to-have.** Strongest argument for prioritising this section. The *endowment effect* explains why the allocation must spend the earned coins: 1,750 coins the player worked for carry

weight; an unearned 0/100 slider carries none. And the *IKEA effect* (we value what we partly assemble) is why the allocation — and choosing the share-card variant — makes the final artifact feel like theirs.

- **Targets = the 8 topics the player just played** (כלכלה וחברה, דת ומדינה, מגדר, ביטחון, מדיני-ביטחוני, אחריות ציבורית, סביבה ואקלים, מי מחליט פה?, ושוויון פנים). The list isn't the NGO editorialising — it's the game's own content, which the player has spent sixteen rounds inside
- **Distribute the 1,750 coins across them.** The currency finally does something and the earning becomes retroactively meaningful. Replay stays meaningful (more coins → more to allocate) with no scarcity or pressure mechanic
- **Free text survives as an optional** אחר — the autonomy argument for unconstrained expression stands; it just doesn't ride on the share card next to the NGO's hashtag
- **The share card names topics, not the free text** — neutral, legible to someone who hasn't played, and no moderation exposure

Why not preset non-topic causes: a curated list of "legitimate priorities" is itself an editorial act — *who decided ביטחון counts and עובדים doesn't?* Using the game's own eight topics sidesteps that entirely.

! Scope flag (VP hat): a real allocation mechanic is **new logic and new state** — **not presentation layer**. That's Roman's territory, landing on a build that already has P1 and P2 queued before mid-September. Options: (a) design + spec now, build after ship; (b) minimal v1 — topics as tappable chips with a simple split, no fine-grained slider; (c) full mechanic in Phase 1 and something named comes off the list. **Decide deliberately, not by drift.** → Decision table #5.

Rejected alternatives: convert 1,750 → 100 points (adds a conversion the player must be taught); retune awards so a full game yields ~100 (cleanest conceptually, but small numbers feel unrewarding to collect).

OPEN Also open: coin totals in captures (175 after topic 1, 325 mid-topic 2) don't reconcile cleanly with the observed award table — either bonus סוגיות contribute or there are awards not yet observed. **Roman: what's the full award table?**

0.4 Bugs and problems found

#	Finding	Severity
1	Confetti fires with the results table and covers it for ~2s — celebration on top of the payload. Also fires at 3/5	High — hides the game's whole point
2	Party is shown before the MK prediction — see §1.4b, this is the root cause of the boring table	High
3	Avatar preview scrolls off screen during customisation — you pick skin tone and hair colour blind	High — trivial fix (position: sticky)
4	Empty white box on the איך היית מצביע? bill card, under the date. Present in stills and video	OPEN Roman: bug or empty state?

#	Finding	Severity
5	Doubled header — avatar and coin count appear twice (top bar + sub-header). At 360px, header + progress card eats ~40% of viewport before content	Medium — fixed by deletion
6	Progress numbers appear inconsistent in play (Lion, 27 Aug). Map says 8/8 while bonus סוגיות remain	OPEN Roman
7	Map→issue transition is a hard cut, no motion	Low — step 7 (juice)
8	Free text goes out verbatim with #המגדלור attached. The content filter is a client-side wordlist — catches slurs, will not catch a party name, a candidate, a slogan, or anything contextual. Arbitrary text carries the NGO's hashtag	High — reputational, not cosmetic
9	דיווח על תוכן פוגעני is an unlabelled red-flag glyph in the corner of the white card. On static hosting there is likely nowhere for a report to go	OPEN Roman — where do reports route?
10	Placeholder promises a mechanic the input doesn't have (\$0.3)	Medium
11	The end-screen white panel is the only surface inverting the dark theme, and isn't using that contrast for anything — reads unfinished rather than deliberate	Low — visual phase

0.5 What's already good — keep

- ↔ החלק את הכרטיס או לחץ למטה — **dual input already shipped**, matching the Tinder/Reigns finding exactly
- הצעה אמיתית בכנסת badge — a credibility marker doing real work
- מה זה? contextual explainer pill — good instinct
- Inline source link (מקור: mako) with a dotted-underline glossary term — **the sentence-completion raw material already exists**
- Map completion states — saturated colour + + הושלם
- Guest skip (כניסה כאורח) — the escape hatch that makes avatar relocation safe
- Free-choice map with no locking — matches §3.1 recommendation already

0.6 Question inventory — read from Tamar's sheet, 27 Aug

8 topics confirmed, ~16 distinct questions (rows 18/19, 20/21, 12/24, 6/16 are duplicates).

KEY Two different kinds of "hard" — the key distinction

Type	What it is	Example	Value
Hard from ignorance	The player doesn't know which bloc an MK sits in	Lion scored 2/5 on — תקציב המדינה a straight 64–55 party-line vote — because he wasn't sure about Mansour Abbas	Fine, but it's a knowledge test in a prediction costume. The reveal teaches; it doesn't surprise
Hard from genuine unpredictability	Someone who knows the Knesset well would still get it wrong	Edelstein voting by mistake; Gantz against his own framework; 99–9	This is the real payload. Roughly 7 exist in the whole sheet

The seven genuine surprises — and where they currently live

#	The surprise	Where it is now
1	99-9 — coalition and opposition voting together on unilateral recognition; Lapid and Gantz with Netanyahu	✓ Already the שקר/אמת question. Strongest in the sheet
2	Gantz voted against his own framework (draft law)	✓ Already the question (row 4 version)
3	Edelstein voted for an opposition amendment by mistake , meaning to abstain — recorded in the protocol	✗ הערות column only
4	Gamliel deliberately absent as silent protest	✗ הערות column only
5	Barkat & Akunis — against in cabinet, for in plenum, to avoid embarrassing the coalition	✗ הערות column only
6	מחנה מכלתי ministers voted against their own government's budget	✗ הערות column only
7	Domestic-violence protection bill fell by a single vote ; Ben-Gvir opposed on false-complaint grounds	✗ תיאור ההצבעה only
—	64-0 because the opposition walked out (reasonableness clause)	✗ Structurally surprising, buried in prose

The curation answer is not "find better questions" — it's "promote what already exists." Five of the seven are already written, verified and sourced, sitting in a notes column the player never sees. Cheapest content win available.

Why these moments matter more than fun — the hypercorrection effect (Metcalfe). Errors made with *high confidence*, once corrected, are remembered better than anything else a person learns. The player who was *sure* Abbas voted for, and discovers he voted against, has just had the strongest learning event cognitive psychology knows how to produce. **The surprise moments aren't just the game's fun — they are literally its most effective teaching.** This is the one-sentence justification of the curation effort, and the line to use with the NGO. Related: the *generation effect* — predicting an answer before seeing it produces far better retention than reading it. The prediction mechanic isn't engagement sugar on civic education; it's the pedagogically stronger delivery mechanism.

Party-line questions — reframed by flow theory (revised v7)

43-48 חוק החמץ · 55-61 חוק בן-גביר · 48-53 ועדת חקירה · חסינות ח"כים · 53-57 הפרדה · 51-65 תקציב המדינה · 55-62 התייעלות · 55-64 היועמ"ש
(4/5) מגדרית

These are not filler — they are the setup. Two principles revise the earlier "cross-bloc everywhere" recommendation: - **Flow (Csikszentmihalyi)**: challenge must sit slightly above skill. With 3 options, chance $\approx 1.7-2.5/5$ — the cross-bloc 2/5 scores are at coin-flip level, meaning *no skill was exercisable*. Prediction that can't beat chance stops feeling like prediction. - **Von Restorff**: a violation only registers against a pattern. If everything is unpredictable, unpredictability is wallpaper.

The rule: pattern, then break. Party-line rounds are where the player *builds* the mental model of the Knesset (and feels competence — 4/5 feels good). Cross-bloc and politics-vote rounds are where the model *breaks* — and per hypercorrection, the break is where the deepest learning lives. Target mix: roughly two-thirds learnable, one-third violating. **Never open the game with a violation.**

! Data problems blocking MK prediction

Row	Problem
נשק אישי	Passed unanimously (— חלקית) — nothing to predict
חוק האקלים	Answer is חלקית — the game only has שקר/אמת buttons. No button exists for this answer
חוק היועמ"ש	התשובה cell is empty
שיתוף פעולה עם הרשות	Attached bill is the Cyber Defence Law — unrelated to the claim
מס כלים · נישואים אזרחיים חד-פעמיים	Marked "סדרו לפי עמדות מתועדות" — no real vote. Credibility rests on documented votes; these need visual distinction (\$0.6b)

PROPOSAL 0.6b Position-based questions need a different treatment

Two questions are built from declared public positions rather than a recorded vote.

The game's entire credibility rests on הצבעות אמיתיות מהכנסת, and the הצעה badge currently appears identically on both kinds. Needs either a distinct badge, or exclusion from MK prediction. → Tamar/NGO.

0.7 data.js findings — read directly, 29 Aug

Verified by reading the shipped data file, not inferred from the build or the question bank. These change build assumptions.

Finding	Consequence
<code>_tally</code> exists on only 8 of 16 issues	The count-up finale — the peak of the round — has structured numbers for half the game . Roman's "vote numbers as fields" question is now blocking , not nice-to-have
Five issues carry no vote numbers anywhere — <code>e2</code> , <code>g1</code> , <code>g2</code> , <code>v2</code> , <code>m1</code>	Not an extraction problem: the numbers don't exist and must be sourced . Tamar ask
<code>v1</code> has <code>tf_answer</code> : "partial" — confirmed in shipped data, not just the sheet	The two-button UI cannot express it. A live bug , not a content question
98 MK cards across 16 rounds, 3–9 per round (<code>r1</code> = 9, <code>v1</code> = 3)	Under predict-then-instant-verdict, 9 cards is a long round. Sharpens the open question: every MK, or <code>key</code> MKs in full and the rest grouped?
Issues carry a <code>core</code> flag — true for the first issue in each topic	A ready-made hook for the $\frac{2}{3}$ pattern / $\frac{1}{3}$ break rule (\$0.6), and for degrading gracefully where a bonus question is missing

Finding	Consequence
22 glossary terms exist and nothing surfaces them	Content already written, invisible. Cheap win — the מה זה? pill is the natural home

1. The reveal — Priority 1

PROPOSAL 1.0 The round arc — revised again, 29 Aug (verified against data.js)

The claim is no longer resolved mid-round. The earlier order revealed the true/false answer at beat 3, before the MK predictions. Reading the shipped data showed that leaks the answers.

1. CLAIM	אמת או שקר → answered, NOT resolved
2. POSITION	איך היית מצביע? → the player as the 121st MK (§1.4c)
3. CONTEXT	the bill → bill_title + bill_date ONLY, neutral framing
4. CASCADE	MK cards one at a time → predict, then instant verdict, each card
5. REVEAL	claim resolves · real tally counts up · resolution · sources

Two open loops now close in a single finale instead of one closing early and one closing late.

WHY IT MOVED — TWO SPOILERS, BOTH IN THE DATA

tf_explain names MKs whose votes are predicted later, in four issues:

Issue	MKs named in tf_explain	key MKs among them
r1 — חוק הגיוס	Gantz, Netanyahu, Lapid	Gantz, Netanyahu
v2	Smotrich, Liberman	both
s1	Ben-Gvir	yes
m2	Lapid	no

r1's explanation states outright that Gantz voted against his own framework. Under the old order that text appeared ~40 seconds before the Gantz card was dealt. Not a hint leak — the answer, printed early, on three of the game's strongest surprise cards.

bill_summary ends with the tally in 9 of 16 issues. For r1, "passed 63 to 57" is arithmetic proof the coalition held: $63 + 57 = 120$, coalition has 64. A player who notices can predict most of those nine cards from the number alone. This spoils the learnable two-thirds rather than the surprise third — but the pattern is what the break breaks against, so losing it costs the break too (\$0.6).

Both are fixed by **sequence**, not by rewriting sixteen texts. `bill_summary` still needs splitting: **context** → **beat 3**, **outcome** → **beat 5**.

KEY CONSEQUENCE: THE PINNED CLAIM BECOMES A DESIGN OBJECT

The player answers at beat 1 and waits until beat 5, so **their answer must stay visible in round chrome across four beat transitions** — otherwise Zeigarnik produces confusion rather than tension.

That makes it the one element continuously on screen for the whole round, and therefore **the game's most load-bearing piece of visual identity** — a banner, poster, sticker or TV chyron, not a progress pip.

! The metaphor is not neutral. A chyron *reports*; a poster *declares*. The second is closer to what the beat actually is (the player's own commitment, held until the truth arrives); the first is safer politically. **Tamar question, not a taste question.**

If the pinned artifact carries the player's **avatar** alongside their answer, the 121st-MK conceit stops being a title and becomes a thing on screen for the whole round.

The held answer still needs consent (Zeigarnik + information-gap): one line of copy (> את התוצאה נגלה בסוף). An unexplained hold reads as a bug, not suspense.

Time budget: ~10s + ~5s + ~5s + (MK cards × ~6s) + ~15s ≈ **60s** ✓ at 3–5 cards. See \$0.7 — some rounds carry 9.

Risks: lopsided votes make the finale flat (mitigate via curation, \$1.4b); same arc ×16 rounds risks staleness (vary MK counts; curated "politics votes" break the pattern).

SET 1.1 The unfinished sentence — confirmed twice

NYT You Draw It • Fibbage — both frame the prompt as an incomplete sentence the reveal completes, source inline in the completed text.

הצביע נגד, יחד עם 12 ח"כים מהקואליציה. [מקור] ... [player predicts] ... Y ח"כ X, בהצבעה על

The build already has explanation prose + inline source on the result screen — this is **restructuring an existing paragraph**, not adding attribution. **NEVER** — the completed sentence editorialises. Verb + count are facts.

SET 1.2 Answer first, truth second

Fibbage • NYT — the chosen answer holds alone for a beat before the truth arrives. Applies to every MK flip and the finale.

SET 1.3 Wrongness has an author — and it isn't the player

Psych ("KASHISH GOT YOU!") — being fooled is an event someone caused. ✓ ח"כ
לא נכון / טעות ✗ — הכנסת הפתיעה אותך הפעם • גולן הפתיע אותך

SET 1.4 Surprise has two sources

Type	Definition	Who defines it	Treatment
Per-player	Gap between this player's prediction and reality	Emerges per player	Default dramatic material — every MK flip dramatizes the gap, framed per \$1.3
Editorial	Votes surprising in themselves — MK voting against own position for tactics	Tamar/NGO, hand-picked — no formula detects "voted cynically"	1-2 curated questions, special beat regardless of the player's answer

Tempo: MK flips fast (~300–500ms — at the *Doherty threshold* where response still feels instant); the finale count-up carries the slow Balatro-weight beat. **Von Restorff effect:** the row that breaks rhythm is the isolated item, and isolated items are what memory keeps. The tempo system is retention engineering, not decoration. **NEVER** — Balatro's escalating jackpot register.

NOW 1.4b Party-before-prediction is the root cause of the boring table

The MK card shows **יהדות התורה · משה גפני** above the prediction buttons. For most Knesset votes, **party is the answer**. The prediction is near-trivial, so the reveal has nothing to reveal.

The table isn't boring because it's a table — it's boring because the guesses were easy. No amount of tempo, colour or cascade fixes a question whose answer is printed above the buttons.

Do not remove the party label — a 17-year-old who doesn't know these names would be coin-flipping, which is worse and less educational.

The fix is curation: choose MK rows where party does not determine the vote. Two or three genuinely unpredictable MKs beat five predictable ones, and the round gets shorter. → **Highest-leverage content task on the project.** Tamar/NGO can start immediately. Decision table #4.

PROPOSAL 1.4c KEY The player's own vote — use the game's own title

איך היית מצביע? is currently collected and discarded. It is the most personal input in the game and it earns nothing, shows nothing, and goes nowhere.

The game is called **הכנסת ה-121**. **The Knesset has 120 members.** The player is the **121st**. That conceit is in the title and completely unused in the build.

Proposal — count the player's vote as the 121st vote:

- **In the round:** the player's own vote appears as a row in the vote reveal, styled as theirs (their avatar, distinct treatment) alongside the MKs
- **In the finale:** **ההצעה עברה 51-62**. עם הקול שלך: 51-63. — usually it changes nothing, and *that is the civic lesson*, delivered as a mechanic rather than a message
- **Across the game:** **הצבעת עם הרוב ב-5 מתוך 8 סוגיות** — a factual alignment record
- **Into the end-game:** the values allocation stops appearing from nowhere. It's introduced by 16 positions the player already took

! HARD GUARDRAIL — compare to the outcome, never to people or parties. - ✓

הצבעת עם הרוב · הצבעת אחרת מהכנסת ב-3 סוגיות — factual, non-labelling - ✗

אתה נוטה ימינה · הצבעת כמו ח"כ · הצבעת כמו הליכוד

The second list is the **Mako Mazpen anti-pattern with extra steps** — a player-to-party match score, exactly what this NGO was founded against. The line is: compare the player to *what happened*, never to *who they resemble*. → Tamar/NGO sign-off before build.

SET 1.4d Sources belong at beat 5, never beat 2 (new guardrail, 29 Aug)

Material served **between the player's opinion and the reveal** is structurally a nudge: the sequencing itself implies their opinion needed correcting. The same links placed **after** the reveal read as "*go check for yourself*."

- ✓ Sources, `bill_summary` outcome, further reading → **beat 5 resolution**
- ✗ Anything explanatory between beat 2 (own vote) and beat 4 — including "helpful" context
- Beat 3 carries **bill_title + bill_date only**. Neutral framing of what was voted on, nothing more

Restated alongside it: **beat 2 is never scored, never rewarded, never compared to a "correct" answer**. Opinion is not prediction. And end-game allocation is across **topics**, never across parties, blocs or positions (§0.3b, §1.4c).

SET 1.5 Crowd data — undecided IF; decided HOW; act NOW on collection

Status: whether the game shows crowd data at all is open and unhurried. Lion's hunch, aligned with research: **if yes, post-reveal only**.

Data collection — **corrected 27 Aug**. The game is not live and has no players, so nothing is currently being lost. The accurate framing: **if** crowd data is ever wanted, capture must exist *before launch*, not after — that is when real answers start arriving. Pre-launch item for Roman (needs a backend; not presentation layer). Not urgent today. → Decision table #3.

Menu when the feature is decided (Tamar/NGO, in writing): (1) never · (2) post-reveal only (*recommended*) · (3) Superbru lock-to-see. **Floor regardless:** no crowd data on or near the values screen, ever. **Why NGO decides:** editorial policy — the NGO owns the public answer to "does the game influence kids?"; the pattern is engaging and will be re-proposed under deadline pressure; a freelance designer must not set editorial policy on a political product for minors. **Contrast for the pitch:** Metaculus shows community pre-commit *on purpose* — aggregation is their product; independent judgment is ours.

SET 1.6 Hard commit gate; incompleteness is never an error

NYT — reveal disabled until complete; unfilled = yellow + escalating instruction copy; zero error states. **Superbru** — lock = two-step with explicit consequence dialog.

OPEN 1.7 Graded scoring

Superbru RESULT / CLOSE / EXACT. Our version = two tiers: outcome prediction (big, held) + per-MK (small, instant). Open: bonus for calling a curated surprise correctly — aligns scoring with the game's thesis. Economy pass, after \$0.3 is settled.

SET 1.8 Describe the shape of the guess

NYT — narrate the pattern before the score: — ניחשת שרוב הקואליציה תצביע יחד
NEVER — characterise the guesser. ובחמישה מתוך שמונה מקרים זה נכון.

OPEN 1.9 Player's guess gets the saturated colour

NYT (guess = hero colour, reality = grey). Cheap to flip; build both. Support with **line/style differentiation**, never colour alone.

SET 1.10 Confetti — cut from the round, keep for the map

Currently fires at 3/5 and **covers the results table for ~2 seconds**. Two failures: it celebrates a coin-flip, and it obscures the payload.

Recommendation: remove confetti from the per-round reveal entirely. Keep it for **map completion (8/8)** only — the one moment it's earned. Scarcity is what makes it mean anything; firing it 16 times makes it wallpaper. (*Peak-End Rule: the game's ending owns the celebration budget — that's where the memory forms.*) If any per-round celebration survives, it lands **after** the table has settled, never over it, and scales to the result.

1.11 Discrepancy as a UI primitive

Papers, Please — **HIGHLIGHT DISCREPANCIES**: a visual language linking two things that don't match. Use for editorial-surprise rows (stated position ↔ actual vote). Screenshots suffice; don't buy.

2. The MK cards

SET 2.1 Reigns: tilt-preview + deliberation as the taught skill

Tilting fades in the answer label **before release**. And its own copy: "**Take a breath! Swipe SLOWLY to see your answers and their consequences.**" The anti-timer, verbatim.

SET 2.2 Tinder: dual input, stamps, calm empty state

Every gesture has a button twin — **already shipped in the build**. Stamp overlay at drag threshold. End-of-stack = calm, one line. **OPEN Playtest question:** does "swipe right = yes" hold in Hebrew RTL or invert? Dual input means the build doesn't block on the answer.

OPEN 2.3 MK representation — spectrum, decided at the visual phase

Currently: coloured circles with Hebrew initials (מג, יא). Functional, but no face means no recognition, and recognising real MKs is itself civic learning.

1. **Raw photos, uniform crop** — credible, zero risk
2. **Treated photos** — duotone/posterise over real likeness: cohesion + reality anchor.
Likely sweet spot
3. **Stylised-realistic redraw** — real likeness in the game's illustration style. Costs: per-MK production across dozens of faces, consistency, and every simplification of a politician's face is readable as intent
4. **~~Caricature~~** — **NEVER**. Mako's bobbleheads are the proof: exaggerating a politician is a political statement, and a neutrality-promising NGO can only lose that game

NOW Open now regardless of level: photo sourcing/rights (Knesset official portraits) — NGO/legal check, not a guess.

3. Topic map — Priority 2

SET 3.1 Free choice; the path is the shape, not the gate

Already free-choice in the build — **keep the logic, change the shape**. Locking topic order would mean "we decide what you learn first" — paternalism from an autonomy-thesis NGO — and it blocks per-topic replay. The ending already exists without locking: celebration gates on 8/8 in any order. Soft touch: a suggested first node for choice-paralysis.

PROPOSAL 3.2 Progress model — 8/8 with split rings (*Lion's proposal, endorsed*)

The build tracks 0/8 topics while each topic contains 2 סוגיות + bonus. Two numbers, one shown.

Recommendation: keep 8/8 as the headline and put sub-progress in the node itself — Duolingo's ring, segmented per סוגיה:

- Ring around each topic node divided into **2 segments** (one per סוגיה)
- 0 filled = untouched · 1 filled = half · 2 filled = הושלם +
- Bonus סוגיות = a **separate small marker** on the node, never part of the completion ring — bonus content must not make הושלם feel like a lie

Why not 0/16 : 16 is a longer, more intimidating number for a "one minute" game, and the topic is the unit players actually choose. The ring gives granularity without changing the headline. **OPEN** Also fixes finding \$0.4-6 — the 8/8 + "bonus remains" contradiction becomes legible.

SET 3.3 Node-state vocabulary (confirmed on camera)

Element	Source
Progress ring on current node (reads without colour)	Duolingo
Locked/done states: grey vs saturated +	Duolingo, Candy Crush, current build
Pinned section header swapping on scroll	Duolingo
Reward/variety nodes mid-path	Duolingo chests
Jump-back-to-current floating button	Duolingo
Dashed connector + numbered nodes + ribbon labels	Monopoly Go world map
Explicit incomplete state — the red "Pick needed" warning	Superbru
Character idling beside the path	Duolingo → OPEN our avatar could live here

NEVER — timer chips, energy, Claim buttons, spin wheels.

OPEN 3.4 Navigation — fully open

מפה / דירוג / ממשלה / אודות : three of four tabs are non-gameplay chrome for a one-minute game. Nothing set in stone, including whether tabs exist. The arc and the map decisions determine what navigation is needed.

4. Character / avatar**SET 4.1 Kill creation-as-first-step**

Currently step 2 of the game, before any play. Research basis: **Duolingo and Mii both start you complete** — Duolingo's taps-to-done is 1.

New placement: default avatar assigned instantly (guest included) → lives small in the map corner, tappable → customisation is **discoverable, not mandatory** → 1-2 contextual invitations at identity moments (after first completed topic; **pre-share**, where the avatar earns its existence).

NOW **Fix regardless of placement:** the preview scrolls off screen during customisation (\$0.4-3). position: **sticky**. **NEVER** — brands, collections, rarity. The avatar is not a store.

5. Share & the end-game**SET 5.1 Mechanism free, artifact is the work**

The build already uses the **native iOS share sheet** — correct, and `navigator.share()` is nearly free. All budget goes to the artifact.

5.2 Current share text — verified 27 Aug

המגדלור # [free text] צברתי 1750 מטבעות, ובחרתי:] " החכ ה-121 " שיחקתי

The bracketed part is whatever the player typed, pasted in verbatim. Confirmed by capture: Cjsjajxjkakxkx passed straight through to the share sheet.

Three weaknesses: 1. **1,750 coins is meaningless to someone who hasn't played** — a number with no scale and no story, and it's the headline 2. **The most personal part of the message is unvalidated user text** carrying the NGO's hashtag (\$0.4-8) 3. Nothing about the *game* is in it — no prediction record, no surprise, no evidence anything happened

PROPOSAL 5.3 KEY What the share should carry instead

" החכ ה-121 " שיחקתי
הכנסת הפתיעה אותי 4 פעמים • ניחשתי נכון 12 מתוך 16
המגדלור # הקצאתי הכי הרבה ל: סביבה ואקלים

KEY Lead with the surprise count, not the prediction record (*Lion, 27 Aug — important correction*).

If the card headlines **ניחשתי נכון 12 מתוך 16**, players who scored 4/16 quietly won't share. That's a self-selection problem twice over: it suppresses sharing, **and** it makes doing badly feel shameful — contradicting the entire §1.3 framing that the Knesset surprised you, you didn't fail.

The metric to lead with is the one that is **high when the player did badly**:

הכנסת הפתיעה אותי 11 פעמים

Someone who scored 4/16 now has the most impressive card in the group. The shame inverts, and it's on-thesis — the game is about the Knesset being surprising, not about the player being right. Same property that makes Wordle's grid shareable at 3/6 and 6/6 alike.

Three moves: - **Lead with the surprise count** — high when you did badly, on-thesis, no self-selection - **Demote the prediction record** to secondary, optional, or a variant card - **Name the topic, from the fixed list of eight** (\$0.3b), not free text. Neutral, legible, no moderation exposure

On the exposure question — revising an earlier over-flag. I had marked sharing the player's priorities as a social-exposure risk. With the eight game topics as the vocabulary, **that flag was too strong**: topic names are subject areas, not positions. **הקצאתי הכי הרבה לסביבה ואקלים** says what you care about without saying what you believe about it.

The narrower residual risk stands: *how you split between topics* maps loosely onto political camps (ביטחון vs. סביבה). So: - **Name the top topic, or the top two** ✓ - **Don't publish the full split** (40% סביבה 10% ביטחון,) unless the player opts into that card - Framing stays מה שחשוב לי , never העמדות שלי - Multiple card variants (Wrapped pattern) — one omits topics entirely, for players who'd rather share only the record

Free text keeps a home as the optional אחר inside the game (\$0.3b). The autonomy argument for unconstrained expression is real — it just doesn't need to travel attached to the NGO's hashtag.

SET 5.4 Wrapped: data as the illustration + player picks the card

Genres-as-sandwich-layers — the data IS the illustration. Player chooses from 3–4 variants; choosing is the self-expression. Bar: survives a teenager's story without embarrassing them.

OPEN 5.5 What data the card carries

Candidate	Verdict
Coins alone	✗ Currently the headline; meaningless to outsiders, and the currency does nothing (\$0.3)
Prediction record (16 מתוך 12 ניחשתי נכון)	❗ Demoted — self-selects against low scorers (\$5.3)
Surprise count (הכנסת הפתיעה אותי 11 פעמים)	✓ ✓ Headline. Neutral, on-thesis, and <i>higher</i> when the player did badly
Top topic named, from the fixed eight	✓ Per \$5.3 — neutral, no moderation exposure
Free text typed by the player	✗ Carries #המגדלור with no real validation (\$0.4-8)
Allocation percentages	❗ Opt-in variant only
Alignment with the Knesset (הצבעתי עם הרוב ב-5 מתוך 8)	OPEN Strong if \$1.4c ships — and safe, because it compares to the outcome, not to parties

Depends on: the success metric (\$10-6). If success = shares, this card is the growth engine and deserves real design days.

6. Counter-references — the מה אנחנו לא בונים slide

Artifact	Shows
Mako Mazpen HaBchirot (Lion's capture)	The tool the NGO's founder criticised publicly. Bobblehead party leaders on podiums, match scores 86/88/90% under המפלגות שהכי מתאימות לך — the whole anti-pattern in one frame
Wahl-O-Mat (Lion's capture)	Ranked party list with %-match bars — the international standard of the same pattern

Artifact	Shows
Metaculus pre-commit community %	Not a villain — a contrast: aggregation is their product, independent judgment is ours

One sentence: **matching tools rank parties by you; we show you reality and let you rank yourself.**

7. Hebrew RTL — rules from the phone sweep

Bit · Max/Cal · trading app · Wolt — four shipping Israeli products, zero exceptions:

Rule	Application
Text anchors right, numbers/status anchor left	MK name right, vote mark left
₪ precedes digits (₪56)	any currency display
Western numerals, LTR digit runs inside RTL flow	vote counts, scores
Key-value: key right, value left	finale stats, detail rows
Semantic colour survives RTL	gain/loss, correct/incorrect

OPEN Mixed-script rows (Hebrew name + Latin abbreviation +) need on-device bidi testing.

8. Consolidated NEVERs

Anti-pattern	Source / basis
Crowd expectation pre-commit	Superbru EXPECTATION bar
Crowd data on the values screen, any position	the founding anti-pattern's cousin
Player-to-party or player-to-MK match in any form	\$1.4c — Mako Mazpen with extra steps
Reactive copy leaking odds pre-commit	Superbru "quite an upset!"
Per-question countdown timer	ESPN clock; Reigns teaches the opposite
Escalating jackpot tally	Balatro register
Dramatizing failure as failure	\$1.3/1.4 — the gap gets drama, reality is the author
Failure copy with the player as subject (טעית)	Psych inversion

Anti-pattern	Source / basis
Caricatured politicians (realistic stylisation is open, \$2.3)	Mako bobbleheads
Celebration covering the payload	\$0.4-1, current build
Publishing the allocation split without opt-in	\$5.3
A share metric that only looks good if you scored well	\$5.3 — self-selects against the players we most want sharing
Explanatory material between beat 2 and beat 4	\$1.4d — sequencing implies the opinion needed correcting
Resolving the claim before the MK cascade	\$1.0 — <code>tf_explain</code> names MKs predicted later; it prints the answer early
Showing <code>bill_summary</code> 's tally at beat 3	\$1.0 — $63+57=120$ is arithmetic proof the coalition held
Unvalidated user text carried on the NGO's hashtag	\$0.4-8 — a client-side wordlist is not moderation
A curated list of "legitimate" causes chosen by the NGO	\$0.3b — the game's own eight topics avoid the editorial act
Rarity / brands / collections on avatar	Bitmoji, Fall Guys
Lock icons on topic content	NYT Games furniture
Timers / energy / Claim / spin on the map	Candy Crush, Monopoly Go
>1 number at the emotional peak	Wordle-clone modal
Characterising the guesser	no political identity from sixteen taps

8b. Principles & evidence — the theory under the decisions *(added v7)*

Evidence tiers: [M] measured (our frame data / playthrough scores) · [O] observed (references or the build) · [L] literature.

Decision	Principle	Evidence
Held outcome + count-up finale	Peak-End Rule — memory keeps the peak and the ending	[L] + [O] NYT holds the gap
Unfinished sentence (§1.1)	Zeigarnik / information gap — an open loop rendered as typography	[O]×2 NYT + Fibbage converged independently
Variable tempo; one row breaks rhythm	Von Restorff — the isolated item is the remembered item	[M] Balatro ~850ms vs Wordle ~125ms
Per-MK instant reveals 300–500ms	Doherty threshold (<400ms feels instant)	[L] + [O] batched table as anti-example

Decision	Principle	Evidence
Prediction before reveal, everywhere	Generation effect — produced answers outlast read ones	[L] — the mechanic is the pedagogy
Curate high-confidence surprises	Hypercorrection — confident errors, corrected, are learning's strongest event	[L] + [M] the 7 moments in \$0.6
Pattern before break; ⅔ learnable	Flow — challenge slightly above skill; chance-level guessing exercises no skill	[M] 2/5 ≈ coin-flip on cross-bloc sets
3 MKs, not 5	Hick's Law — fewer, meaningful choices	[M] 4/5 homogeneous set = two dead rows
8/8 headline + ring sub-progress	Miller / chunking — 8 at the edge of working memory, 16 past it	[L]
Party label stays	Tesler's Law — hiding it dumps the complexity on the player as noise	[L]
Swipe + tap everywhere	Postel's Law + hedges the RTL unknown	[O] shipped
Familiar atoms, novel sequence	Jakob's Law — users import expectations from other products; the whole reference sprint was this	[O] Tinder/Duolingo/native share/RTL sweep
Ring progress, עוד 7 וסיימת , bonus off the ring	Goal-Gradient — visible receding goals break the pull	[O] the build's own 8/8-with-bonus contradiction
Allocate the <i>earned</i> coins	Endowment effect — worked-for currency carries weight	[L]
Player assembles allocation + picks share card	IKEA effect — we value what we partly built	[L] + [O] Wrapped card choice
Autonomy / competence / relatedness throughout	Self-Determination Theory — the named answer to "engaging without dark patterns"	[L] — cite when asked why no streaks

Guardrails, grounded:

NEVER	Mechanism it blocks
Crowd data pre-commit	Anchoring + conformity — exposed majority opinion drags judgment toward it
Streaks / loss framing	Loss aversion (~2× gains) — streaks weaponise it
Variable-ratio rewards	The Skinner-box schedule — the slot-machine compulsion mechanism; take Balatro's tally, refuse its schedule
Prediction record as share headline	Self-selection under social comparison — low scorers won't post
Political-position sharing on by default	Default effects — the default is the strongest choice-architecture lever; on a product for minors it must be the safe artifact

Deliberately not used: Bartle types (audience not segmentable at this scale) · Hook Model (we are explicitly not building habit loops) · Octalysis (framework overhead, no new insight) · aesthetic-usability (true, but a Gate 1 argument, not a research finding).

9. Method notes

Trim recordings before zipping · Claude can't watch YouTube, files only · 60fps capture + frame measurement when timing matters · Mobbin MCP requires a paid plan, skipped.

10. Decisions needed — routed

#	Question	Owner	Status
1	Round arc adoption (\$1.0) — concept sign-off	Tamar/ NGO	Sketch → Sunday call
2	_tally on only 8 of 16 issues — extract vote numbers into structured fields. Now blocking the finale , not optional	Roman	NOW Blocking
2b	Split <code>bill_summary</code> : context → beat 3, outcome → beat 5	Roman + Lion	With the arc build
2c	<code>v1 tf_answer: "partial"</code> — two-button UI can't express it. Live bug	Roman + Tamar	NOW Quick
3	Start collecting anonymised answer data	Roman	NOW Has a clock on it
4	Question curation — \$0.6: promote the 5 buried surprises; fix the 4 data problems; position-based questions; pattern-then-break (⅓ learnable, ⅓ violating, never open on a violation). The <code>core</code> flag is a ready-made hook	Tamar/ NGO	NOW Highest-leverage content task
4b	Source vote numbers for 5 issues — <code>e2</code> , <code>g1</code> , <code>g2</code> , <code>v2</code> , <code>m1</code> have none anywhere	Tamar/ NGO	NOW Blocks the finale on those rounds
4c	Pinned-claim metaphor : banner / poster / sticker / chyron. A chyron reports, a poster declares — not a taste question	Tamar/ NGO	With Gate 1
4d	MK card count: 3–9 per round. Every MK, or <code>key</code> in full and the rest grouped?	Lion + Tamar	Affects round length
5	End-game: no allocation mechanic exists (\$0.3). Build topic allocation? Which scope option — spec-now-build-later / minimal chips / full in Phase 1? Full award table? Is the DOM 0/100 dead code or half-built?	Lion + Roman	NOW Blocks the end-game; not presentation layer
5b	Where do דיווח על תוכן פוגעני reports route on static hosting?	Roman	NOW Quick answer
5c	Free text on the share card — keep, restrict to in-game אחר, or remove? Editorial call given the hashtag exposure	Tamar/ NGO	Sunday call

#	Question	Owner	Status
6	Success metric: completions, shares, or replays	NGO	Asked, unanswered
7	Player's vote as the 121st (\$1.4c) + the never-compare-to-parties guardrail	Tamar/ NGO	With #1
8	Share-card content policy (\$5.3, 5.5)	Tamar/ NGO	After #6
9	MK photo rights/sourcing	Tamar/ NGO (+legal)	NOW Before MK card build
10	Progress model: 8/8 + split rings (\$3.2)	Lion + Roman	Low friction
11	Avatar placement change vs current build	Lion + Roman	Boundary list
12	RTL swipe direction	Playtest (teens)	Dual input unblocks build
13	Empty white box on איך היית מצביע? — bug?	Roman	NOW Quick answer

11. Immediate build fixes (no sign-off needed)

Presentation-layer, uncontroversial, cheap: 1. Sticky avatar preview during customisation 2. Remove the doubled header — one avatar, one coin counter 3. Confetti out of the per-round reveal 4. Header/progress compression at 360px

Not in this list: the end-game allocation (\$0.3b) — new logic and state, Roman's territory, needs the scope decision first.

12. Parking lot — visual direction (deferred)

- **NYT Games card language:** illustration breaking the card frame • one flat saturated colour per card • card stack with peek (**OPEN** topic-map alternative) • byline-as-attribution → source as design element • "between Gen Z and serious"
- **Balatro:** flat/chunky/high-contrast, adjacent to the codebase's existing neo-brutalist layer
- **Current build:** dark paper, offset shadows, per-topic hues, `SimplerPro_HLAR-Black`. Gate 1 question: keep / sharpen / replace